

Spectral Clustering


Part 1: The Graph Laplacian

Ng Yen Kaow

Laplacian of a function

□ Given a multivariate function $f: \mathbb{R}^n \rightarrow \mathbb{R}$

□ $\nabla f(\mathbf{x})$, the gradient at $f(\mathbf{x})$, is a vector pointing at the steepest ascent of $f(\mathbf{x})$

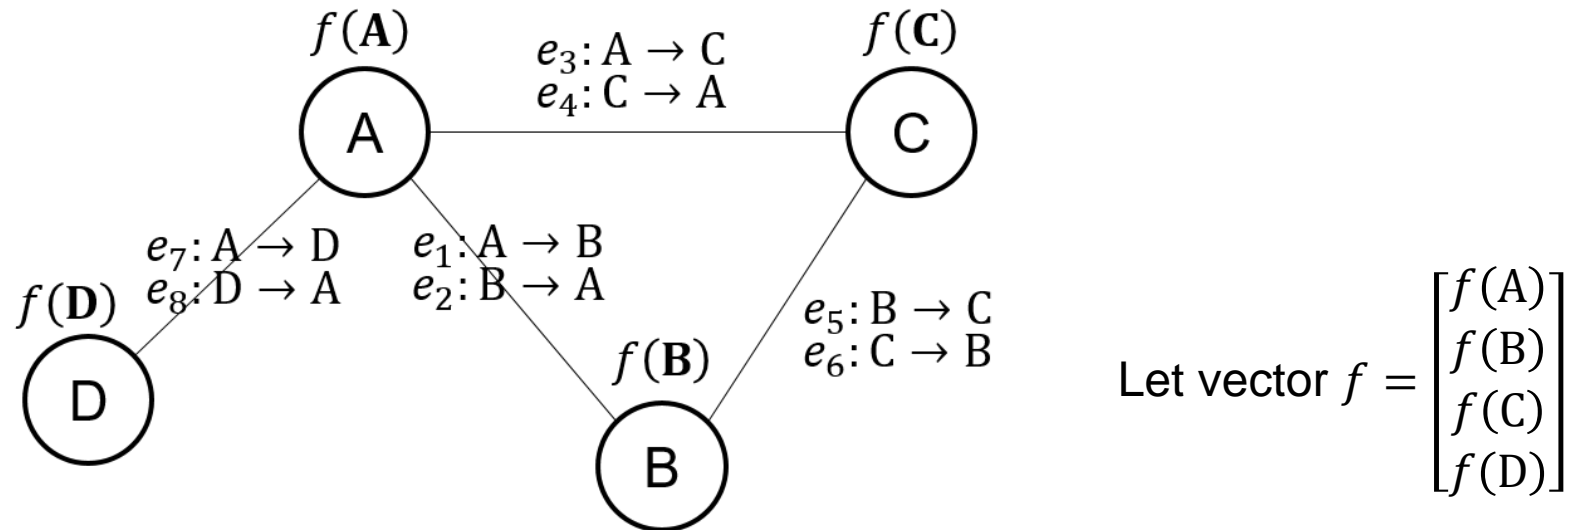


Vector field ∇f

□ Δf , the Laplacian of f , is the divergence of ∇f , that is, $\Delta f(\mathbf{x}) = \nabla \cdot \nabla f(\mathbf{x})$


■ A scalar measurement of the smoothness in $\nabla f(\mathbf{x})$ about point \mathbf{x}

Incidence matrix



- Consider each vertex as a point on the grid
 - The domain of f are now the vertices
 - $f(v)$ operates on each vertex v
 - The gradient from vertex v to v' is given by the edge $e: v \rightarrow v'$, more specifically, $f(v') - f(v)$
 - Denote the gradient of edge e as $w(e)$
- Define a matrix which captures all the gradients

Incidence matrix



Let vector $f = \begin{bmatrix} f(A) \\ f(B) \\ f(C) \\ f(D) \end{bmatrix}$

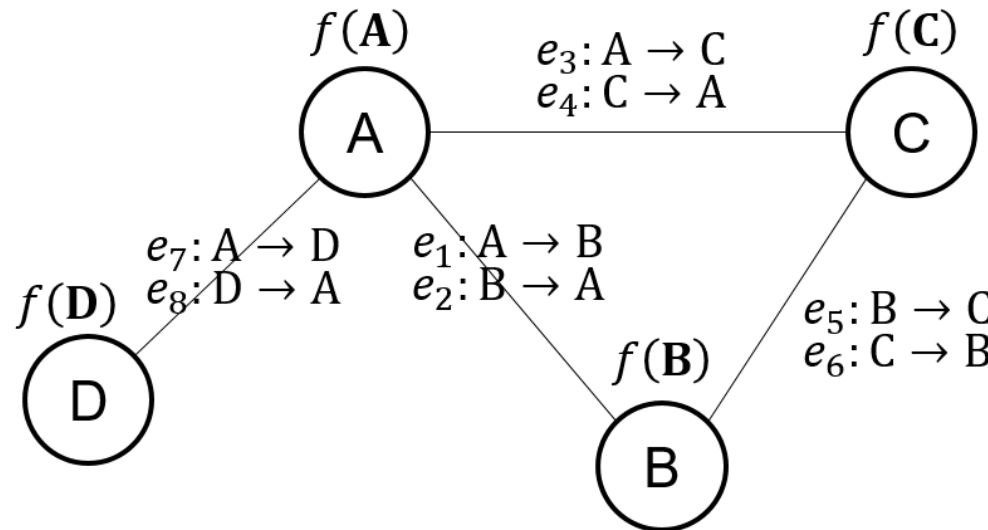
□ Incidence matrix

$$M = \begin{matrix} & e_1 & e_2 & e_3 & e_4 & e_5 & e_6 & e_7 & e_8 \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{bmatrix} 1 & -1 & 1 & -1 & 0 & 0 & 1 & -1 \\ -1 & 1 & 0 & 0 & 1 & -1 & 0 & 0 \\ 0 & 0 & -1 & 1 & -1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & -1 & 1 \end{bmatrix} \end{matrix}$$

□ Every column represents an edge in the graph

$$(M^T)_{\underbrace{1}_{\text{row 1 of } M}} f = [1 \quad -1 \quad 0 \quad 0] \begin{bmatrix} f(A) \\ f(B) \\ f(C) \\ f(D) \end{bmatrix} = f(A) - f(B) = w(e_1)$$

Incidence matrix



Let vector $f = \begin{bmatrix} f(A) \\ f(B) \\ f(C) \\ f(D) \end{bmatrix}$

□ Incidence matrix

$$M = \begin{matrix} & \begin{matrix} e_1 & e_2 & e_3 & e_4 & e_5 & e_6 & e_7 & e_8 \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \end{matrix} & \begin{bmatrix} 1 & -1 & 1 & -1 & 0 & 0 & 1 & -1 \\ -1 & 1 & 0 & 0 & 1 & -1 & 0 & 0 \\ 0 & 0 & -1 & 1 & -1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & -1 & 1 \end{bmatrix} \end{matrix}$$

- Every column represents an edge in the graph
- $M^T f$ is a $|E| \times 1$ vector where each entry gives the gradient of an edge
 - $M^T f$ contains all the gradients of the graph

The graph Laplacian L

- The graph Laplacian L is obtained by

$$\Delta f = \nabla \cdot \nabla f = MM^T f$$

e.g.

$$(MM^T f)_1 = [1 \quad -1 \quad 1 \quad -1 \quad 0 \quad 0 \quad 1 \quad -1] \begin{bmatrix} w(e_1) \\ w(e_2) \\ w(e_3) \\ w(e_4) \\ w(e_5) \\ w(e_6) \\ w(e_7) \\ w(e_8) \end{bmatrix} = \underbrace{w(e_1) - w(e_2) + w(e_3) - w(e_4) + w(e_7) - w(e_8)}_{\text{divergence of vertex A}}$$

- $MM^T f$ is a $|V| \times 1$ vector where each entry gives the divergence of a vertex
- MM^T is a $|V| \times |V|$ matrix where

$$MM^T \begin{bmatrix} f(A) \\ f(B) \\ \vdots \end{bmatrix} = \begin{bmatrix} \Delta f(A) \\ \Delta f(B) \\ \vdots \end{bmatrix}$$

Properties of L

- The graph Laplacian L is obtained as $L = MM^T$
 - Since L is of the form MM^T , L is **symmetric** and **positive-semidefinite**
 - This allows us to obtain an orthogonal eigenbasis, which has special meanings (next slide)
 - $L = D - A$, where D is the degree matrix and A the adjacency matrix

Eigenvectors of L

- The eigenvectors of L has special meaning
 - Consider the vectors x fulfilling $Lx = \lambda x$,
 - Compared with $Lx = \begin{bmatrix} \text{divergence of } A \\ \vdots \\ \text{divergence of } A \end{bmatrix}$, we have that $\lambda x = \begin{bmatrix} \text{divergence of } A \\ \vdots \\ \text{divergence of } A \end{bmatrix}$
 - The eigenvector x corresponds to the values $f(A), f(B), \dots$, where $\lambda f(v) \approx \Delta f(v)$
 - A small λ indicates that $f(v)$ does not vary much from $f(v')$ of its neighbors v'
 - A connected graph has $\min(\lambda) = 0$, indicating that $\Delta f(v) = 0$, (i.e. $f(v) = \text{const}$, a stationary state)
 - For a disconnected graph, the disconnected components has different constants for $f(v)$ values

Mathematical property of L

- A precise mathematical property of L relates it to “sparsest cut” problems
- Let the adjacency matrix $A = (a_{ij})$, then

$$x^\top Lx = \frac{1}{2} \sum_{i,j=1}^m a_{ij} (x_i - x_j)^2$$

$$\begin{aligned} x^\top Lx &= x^\top Dx - x^\top Ax = \sum_{i=1}^m d_i x_i^2 - \sum_{i,j=1}^m a_{ij} x_i x_j \\ &= \frac{1}{2} \left(\sum_{i=1}^m d_i x_i^2 - 2 \sum_{i,j=1}^m a_{ij} x_i x_j + \sum_{i=1}^m d_i x_i^2 \right) \\ &= \frac{1}{2} \sum_{i,j=1}^m a_{ij} (x_i - x_j)^2 \end{aligned}$$

Mathematical property of L

- A precise mathematical property of L relates it to “sparsest cut” problems

- Let the adjacency matrix $A = (a_{ij})$, then

$$x^\top Lx = \frac{1}{2} \sum_{i,j=1}^m a_{ij} (x_i - x_j)^2$$

- Suppose x is a vector of only the values +1 and -1, indicating the membership of the vertices in a set S

$$x_i = \begin{cases} 1 & \text{if } v_i \in S \\ -1 & \text{if } v_i \in \bar{S} \end{cases}$$

- That is, we want to use x to indicate the result of a 2-partition, S and \bar{S}

Mathematical property of L

- A precise mathematical property of L relates it to “sparsest cut” problems
- Let the adjacency matrix $A = (a_{ij})$, then

$$x^{\top} L x = \frac{1}{2} \sum_{i,j=1}^m a_{ij} (x_i - x_j)^2$$

- Suppose x is a vector of only $\{1, -1\}$, then $x^{\top} L x$ has special significance

$$\begin{aligned} \frac{1}{2} \sum_{i,j=1}^m a_{ij} (x_i - x_j)^2 &= \sum_{i,j=1, i < j}^m a_{ij} (x_i - x_j)^2 \\ &= 4 \sum_{1 \leq i < j \leq m, x_i \neq x_j} a_{ij} \end{aligned}$$

- That is, $x^{\top} L x$ is 4 times the number of edges between adjacent vertices of each from S and \bar{S}

Finding x that minimizes $x^T L x$

- Compute $x^T L x$ for all x
 - e.g. $x = [1, -1, -1, -1]$ gives $x^T L x = 12$
- This gives us the 2-partition that results in the least number of removed edges
 - $x = \mathbf{1} = [1 \ 1 \ 1 \ 1]$ or $x = -\mathbf{1} = [-1 \ -1 \ -1 \ -1]$ which has $x^T L x = 0$ are trivial solutions
 - Best x is $[1 \ 1 \ 1 \ -1]$, that is, A, B, C in one group and D in another

Group 1	Group 2	$x^T L x$
A	B C D	12
B	A C D	8
C	A B D	8
D	A B C	4
A B	C D	12
A C	B D	12
A D	B C	8
A B C D	\emptyset	0

- The optimal x can be approximately found

Finding x that minimizes $x^\top Lx$

□ Minimize $x^\top Lx$

■ Consider instead problem of minimizing $\frac{x^\top Lx}{x^\top x}$

□ x is of only +1 and -1 $\Rightarrow x^\top x = |x| = \text{const}$

Group 1	Group 2	$x^\top Lx$	$\frac{x^\top Lx}{x^\top x}$
A	B C D	12	3
B	A C D	8	2
C	A B D	8	2
D	A B C	4	1
A B	C D	12	3
A C	B D	12	3
A D	B C	8	2

Finding x that minimizes $x^\top Lx$

□ $\frac{x^\top Lx}{x^\top x}$ is known as the Rayleigh quotient

■ By the min-max theorem of Rayleigh quotient,

$$\min_x \frac{x^\top Lx}{x^\top x} = \lambda_k$$

■ where λ_k is the smallest eigenvalue in the decomposition of $Lx = \lambda x$, and

■ $\mu_k = \operatorname{argmin}_x \frac{x^\top Lx}{x^\top x}$

□ However, μ_k is the trivial ($\lambda_k = 0$) solution

■ Compromise and use the second best solution μ_{k-1} (which corresponds to the second smallest eigenvalue λ_{k-1})

Eigendecomposition example

□ Eigenvalues

λ_1	λ_2	λ_3	λ_4
4.0000	3.0000	1.0000	0.0000

□ Eigenvectors

μ_1	μ_2	μ_3	μ_4
0.8660	0.0000	0.0000	-0.5000
-0.2887	0.7071	-0.4082	-0.5000
-0.2887	-0.7071	-0.4082	-0.5000
-0.2887	0.0000	0.8165	-0.5000

More precisely, -9.51E-17

- $\lambda_3 = 1 = \text{optimal value for } \frac{1}{2} \sum_{1 \leq i, j \leq m} a_{ij} (x_i - x_j)^2$
- If group by the (\pm) sign, μ_3 correctly places A, B, C in one group ($-$) and D in another ($+$)

Compromise in +1/-1 restriction

- By relaxing the restriction of +1 and -1 in x to allow any real number, an $x^T L x$ smaller than the optimal under the restriction is often achieved
- The improvement can be guaranteed if x is orthogonal to $\mathbf{1}$ (or $-\mathbf{1}$) since by the min-max theorem, $\frac{\mu_{k-1}^T L \mu_{k-1}}{\mu_{k-1}^T \mu_{k-1}}$ is minimal among all $\frac{x^T L x}{x^T x}$ that are orthogonal to μ_k
 - However, in the present case, $x = [1 \ 1 \ 1 \ -1]$ and not orthogonal to $\mu_4 = [1 \ 1 \ 1 \ 1]$
 - Still, $\frac{\mu_3^T L \mu_3}{\mu_3^T \mu_3} = \lambda_3 = 1 = \min_{x \in \{1, -1\}^4} \frac{x^T L x}{x^T x}$
 - Though no guarantee, improvements are usual

The significance of μ_{k-1} and λ_{k-1}

- The heuristic for translating μ_{k-1} back into discrete values for a grouping of the vertices is an important topic
- μ_{k-1} is called the **Fiedler vector**
- λ_{k-1} is called the **Fiedler value**
 - The multiplicity of λ_{k-1} is always 1
 - Also called the **algebraic connectivity**
 - The further λ_{k-1} is from 0, the more connected is the graph